

---

# Generative Pretraining from Pixels

Mark Chen, Alec Radford, Rewon Child, Jeff Wu, Heewoo Jun, Prafulla Dhariwal,  
David Luan, Ilya Sutskever

ICML 2020

Presented by Mingyu Yang

---

# Background

---

- **Unsupervised generative pre-training for images:**
  - Popular in mid 2000's
  - A central role in the resurgence of deep learning:
    - Before that, DNNs are very hard to train!
    - People believe that learning  $p(x)$  helps supervised modeling of  $p(y|x)$
    - Pre-training + fine-tuning achieves STOA performance and outperforms SVM in MNIST


# Background

- **Unsupervised generative pre-training for images:**

- Popular in mid 2000's
- A central role in the resurgence of deep learning:
  - Before that, DNNs are very hard to train!
  - People believe that learning  $p(x)$  helps supervised modeling of  $p(y|x)$
  - Pre-training + fine-tuning achieves STOA performance and outperforms SVM in MNIST
- Example:
  - Deep Belief Network (2006)
  - Denoising Autoencoder (2008)



Digits generated from Deep Belief Networks



Manifold learning perspective of denoising autoencoder

# Background

---

- Unsupervised generative pre-training becomes less popular for images:

1. Deep Neural Networks are much easier to train
  - Better activation functions: *ReLU*, *LeakyReLU*, ...
  - Improved initializations: *Xavier*, *Kaiming*, ...
  - Normalization strategies: *Batch Normalization* ...

# Background

---

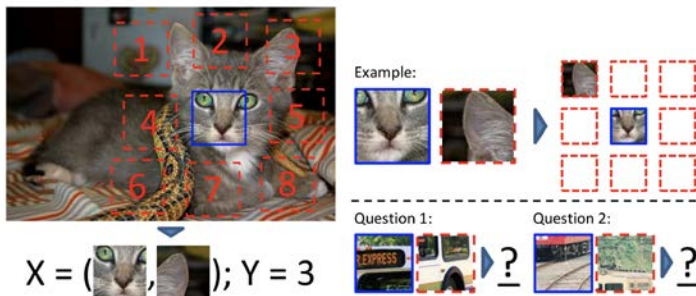
- Unsupervised generative pre-training becomes less popular for images:
  1. Deep Neural Networks are much easier to train
    - Better activation functions: *ReLU, LeakyReLU, ...*
    - Improved initializations: *Xavier, Kaiming, ...*
    - Normalization strategies: *Batch Normalization ...*
  2. Supervised pre-training achieves better performance

# Background

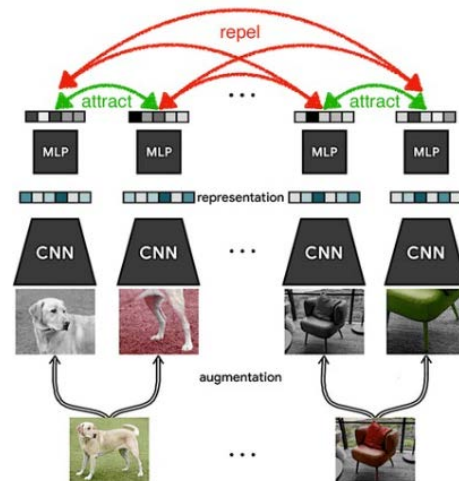
- Unsupervised generative pre-training becomes less popular for images:

3. People design different self-supervised learning methods model global structures (by solving pretext tasks) instead of the distribution

- Predict relative positions
- Mutual Information: *AMDIM*
- Contrastive learning: *MoCo*, *SimCLR*



Predicting relative positions (Doersch et al., 2015)



# Background

- **Unsupervised generative pre-training flourished in NLP!**

- Learn the language model as pre-training

- BERT (2018):  $L_{BERT} = \mathbb{E}_{x \sim X} \mathbb{E}_M \sum_{i \in M} [-\log p(x_i | x_{[1,n] \setminus M})]$

Predict masked words: **Conditional probability**

- GPT-2 (2019), GPT-3 (2020):  $L_{AR} = \mathbb{E}_{x \sim X} [-\log p(x)]$


$$p(x) = \prod_{i=1}^n p(x_{\pi_i} | x_{\pi_1}, \dots, x_{\pi_{i-1}}, \theta)$$

**Autoregressive model**


# Background

- Unsupervised generative pre-training *flourished in NLP!*

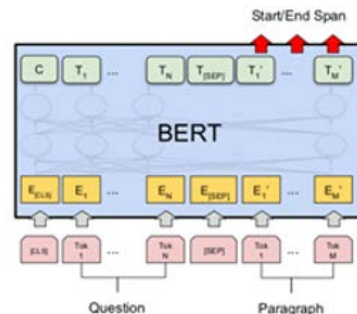
- Fine-tuning for different tasks




(a) Sentence Pair Classification Tasks:  
MNLI, QQP, QNLI, STS-B, MRPC,  
RTE, SWAG



(b) Single Sentence Classification Tasks:  
SST-2, CoLA



(c) Question Answering Tasks:  
SQuAD v1.1



(d) Single Sentence Tagging Tasks:  
CoNLL-2003 NER



# Background

- **Unsupervised generative pre-training flourished in NLP!**

- STOA methods benefit from attention mechanism


- Transformer (2017)
- Will be introduced later in this course

- Motivation of this paper:

- Can we do the same to images?
- Can we get competitive performance?

- Very similar to BigBiGAN presented last class


- Autoregressive vs GAN



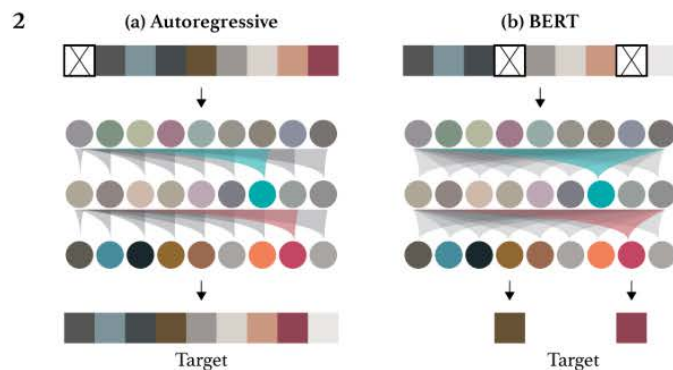
Transformer architecture

# Method


- Overview



Pre-processing



Pre-training




Fine-tuning

# Method

- **Pre-processing:**

- Context Reduction:
  - Images are so large for transformers
    - ImageNet: 224x224x3
  - Downsampling:
    - Reduce the size to 32x32x3, 48x48x3, or 64x64x3
  - Reduce 3 dimensional (R,G,B) channels to 1 dimensional using K-means
    - 512 clusters
    - Further reduce the size to 32x32, 48x48, or 64x64




# Method

- **Pre-training:**

- Autoregressive (AR):

- Likelihood:  $p(x) = \prod_{i=1}^n p(x_{\pi_i} | x_{\pi_1}, \dots, x_{\pi_{i-1}}, \theta)$
    - Minimize log-likelihood:  $L_{AR} = \mathbb{E}_{x \sim X} [-\log p(x)]$
    - Raster order
    - Upper triangular mask to zero out the effect of future words (pixels)




# Method

- **Pre-training:**

- Autoregressive (AR):

- Likelihood:  $p(x) = \prod_{i=1}^n p(x_{\pi_i} | x_{\pi_1}, \dots, x_{\pi_{i-1}}, \theta)$
    - Minimize log-likelihood:  $L_{AR} = \mathbb{E}_{x \sim X} [-\log p(x)]$
    - Raster order
    - Upper triangular mask to zero out the effect of future words (pixels)




# Method

- **Pre-training:**

- BERT:

- Minimize the negative log-likelihood of the masked elements conditioned on the unmasked ones

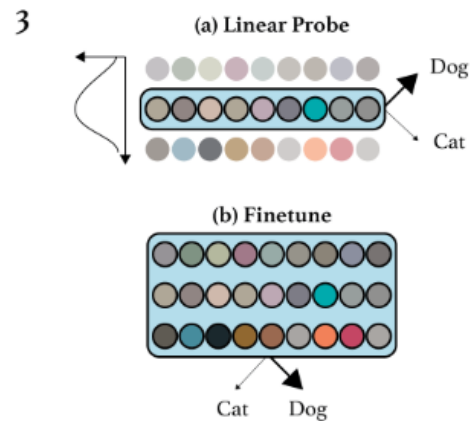
$$L_{BERT} = \mathbb{E}_{x \sim X} \mathbb{E}_M \sum_{i \in M} [-\log p(x_i | x_{[1,n] \setminus M})]$$



# Method

- **Fine-tuning:**

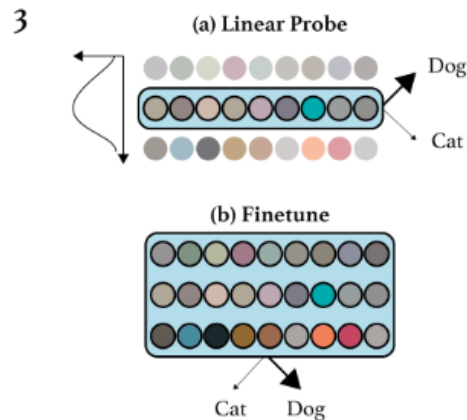
- At each layer, each pixel gets a *d-dimensional* feature vector
- *Average pooling* across the sequence dimension to get a *d-dimensional* feature vector for the whole image
- Linear Probe & Fine-tuning



# Method

- **Fine-tuning:**

- Linear Probe (transfer learning):
  - Treat the transformer as a fixed feature extractor
  - Learn a projection to class logits and minimize the cross entropy loss  $L_{CLF}$
  - Could extract the features at any intermediate layer

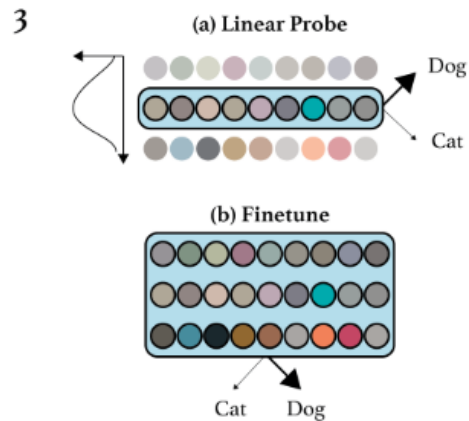




# Method

- **Fine-tuning:**

- Fine-tuning:
  - Treat the learned transformer as an initialization
  - Learn a projection to class logits and minimize the joint objective  $L_{GEN} + L_{CLF}$



# Method


- **More details:**
  - Same model with GPT-2 with slight modification

	# of layers	Embedding size	# of parameters
iGPT-S	24	512	76M
iGPT-M	36	1024	455M
iGPT-L	48	1536	1.4B
iGPT-XL	60	3072	6.8B

Model parameters


# Evaluations for AR

- **Representations at different layers:**
  - Linear probes for the features at different layers
  - Middle layers works the best (Not the last layer?)



# Evaluations for AR

- **Better generative models learn better representations:**
  - Dotted markers denote checkpoints at steps 65K, 131K, 262K, 514K, and 1000K



# Evaluations for AR

- **Linear probes on CIFAR and STL-10:**
  - Pre-trained on ImageNet
  - Outperform STOA unsupervised pre-training methods

- pre-train with 32x32 down-sampled images
- linear probes with 32x32 images

- pre-train with 224x224 images
- Transfer learning with 224x224 up-sampled images

Model	Acc	Unsup Transfer	Sup Transfer
<b>CIFAR-10</b>			
ResNet-152	94		✓
SimCLR	95.3	✓	
iGPT-L	96.3	✓	
<b>CIFAR-100</b>			
ResNet-152	78.0		✓
SimCLR	80.2	✓	
iGPT-L	82.8	✓	
<b>STL-10</b>			
AMDIM-L	94.2	✓	
iGPT-L	95.5	✓	

# Evaluations for AR

- **Linear probes on ImageNet:**
  - Comparable but not better performance than STOA (at that time)
  - Much larger model and longer features

Method	IR	Params (M)	Features	Acc
Rotation	orig.	86	8192	55.4
iGPT-L	$32^2 \cdot 3$	1362	1536	60.3
BigBiGAN	orig.	86	8192	61.3
iGPT-L	$48^2 \cdot 3$	1362	1536	65.2
AMDIM	orig.	626	8192	68.1
MoCo	orig.	375	8192	68.6
iGPT-XL	$64^2 \cdot 3$	6801	3072	68.7
SimCLR	orig.	24	2048	69.3
CPC v2	orig.	303	8192	71.5
iGPT-XL	$64^2 \cdot 3$	6801	15360	72.0
SimCLR	orig.	375	8192	76.5

# Evaluations for AR


- **Full fine-tuning:**

- Pre-trained on ImageNet
- Achieve STOA performance on CIFAR-10
- Achieve 66.3% on ImageNet with 32x32x3, which is worse than the STOA performance of 70.2% (Isometrix Neural Nets)

Model	Acc	Unsup Transfer	Sup Transfer
<b>CIFAR-10</b>			
AutoAugment	98.5		
SimCLR	98.6	✓	
GPipe	99.0		✓
iGPT-L	99.0	✓	
<b>CIFAR-100</b>			
iGPT-L	88.5	✓	
SimCLR	89.0	✓	
AutoAugment	89.3		
EfficientNet	91.7		✓

# Evaluations for AR and BERT

- **BERT vs AR:**



- AR outperforms BERT with linear probes
- BERT catches up with fine-tuning



# Summary

---

- **Pros:**

- Transformers do work for image pre-training
- Achieves impressive performance on small dataset such as CIFAR-10
- Large potential for future improvements (e.g., combine with CNN for large images, GPT-3, etc)

- **Cons:**

- Difficult to deal with high resolution images. Downsampling causes a loss of information
  - Huge memory and computation cost
  - Ignoring the spatial information
-